

AI EXP : HW3 : Linear Regression  
Alex Lamarche

**Problem 1: Setting up the functions**

1.A - Scripts included or check the GitHub link here : [https://github.com/ajl3545/AI\\_HW3](https://github.com/ajl3545/AI_HW3)

1.B – Derivation of  $g(w)$  the gradient function

$$\begin{aligned}
 C(w; X, y) &= L(w; X, y) + \lambda R(w) \\
 \nabla C(w; X, y) &\rightarrow \nabla_w \sum_{n=1}^N |y_n - h(w; x_n)|^2 + \|w\|_2^2 \\
 A_n &= \begin{bmatrix} 1 & x_n \end{bmatrix} \text{ and,} \\
 \nabla L(w; x, y) &= \nabla \|A_{tr} w - y_{tr}\|_2^2 \\
 &= \nabla_w (A_{tr} w - y_{tr})(A_{tr} w - y_{tr}) \\
 &= \nabla_w A_{tr}^T A_{tr} w - 2 A_{tr}^T y_{tr} \\
 &= \nabla_w A^T A w - \nabla_w 2 A_{tr}^T y_{tr} \\
 &= A^T \nabla_w w^2 - 2 A_{tr}^T y_{tr} \nabla_w w \\
 &= 2 w A^T - (2 A_{tr}^T y_{tr}) \cdot 1 \\
 &= \underbrace{(2 w^T A^T A - 2 A_{tr}^T y_{tr})}_{g(w)}
 \end{aligned}$$

1.C Derivation of  $g(W_{opt})$  the gradient function whose value is:  $g(w) = 0$

$$\begin{aligned}
 0 &= 2 w^T A^T A - 2 A_{tr}^T y_{tr} \\
 2 A_{tr}^T y_{tr} &= 2 w^T A^T A \\
 w_{opt} &= \frac{A_{tr}^T y_{tr}}{A^T A} \\
 w_{opt} &= (A^T A)^{-1} A_{tr}^T y_{tr} \\
 \text{with regularization so} \\
 w_{opt} &= (A^T A + \lambda I)^{-1} A_{tr}^T y_{tr}
 \end{aligned}$$

## 1.D – Gradient Algorithm presentation and effects of $\alpha$ and $\tau$ :

```
def GD_SOLVER(X, y, p, l, step):

    parameters = []; parameters.append(p)
    costs = []

    i = 1

    while (True):

        # Most recent W
        w = parameters[-1]

        (mtr) = REG_MET(X, y, w, l)
        costs.append(mtr)

        g = gradient(X,w,y,l)

        # Terminate
        if (np.linalg.norm(g)**2 < pow(10,-8)):
            print("W len = " + str(len(parameters)))
            print("iters len = " + str(i))
            return (parameters, costs)

        # Descend gradient
        parameters.append(w - step*g)

        i+=1

def gradient(X,w,y,l):
    A = []
    for row in X:
        A.append(np.append(1,row))
    A_tr = np.array(A)
    A_trans = np.transpose(A_tr)

    d = np.dot(A_trans,A_tr)
    ident = l*np.identity(len(d))

    return (2 * np.dot(d+ident,w)) - (2 * np.dot(A_trans,y))
```

*Discussion* –  $\alpha$  represents the step that must be taken towards the minimal value of the computed gradient. Each of the  $w$  values changes with the gradient descent iterations - towards an optimal set of parameters. By subtracting the previous set of  $w$  parameters by the current gradient “direction,” computed by  $\text{gradient}(W,x,y,z)$ , the algorithm gets closer to a solution that minimizes the cost function. Step is a multiplier that is used to get closer to the minimum of the gradient function. Increasing the step could lead to overfitting and would overshoot on the gradient.

$\tau$ , a threshold per se, is used to compute how close the parameters get us to an optimal regression line. First, we step towards the minimum, then we check how close we are.  $\tau$  ensures that we are close to the minimal value. If the summation of the squares of the parameters (the cost at that parameter iteration) is lower than  $\tau$ , then that means the most recent  $w$  parameters have- as closely as possible and within reason – calculated a minimal cost.

\*ignore the print statements\*

## Problem 2 – Linear Regression (true function is affine)

2.A [https://github.com/ajl3545/AI\\_HW3/blob/main/problem2.py](https://github.com/ajl3545/AI_HW3/blob/main/problem2.py) or check out included file

2.B – Present the results of CF\_SOLVER on training and testing data and show MSE for each:

Training data:

```
Wopt =
[0.9999999999999999, 2.0000000000000001, 2.9999999999999999, 4.0, 4.9999999999999999, 5.0, 3.9999999999999999, 2.9999999999999999, 1.9999999999999998, 1.0000000000000002]
Opt Cost (mtr) = 3.3695454009416875e-28
Opt MSE = 3.7439393343796526e-29
```

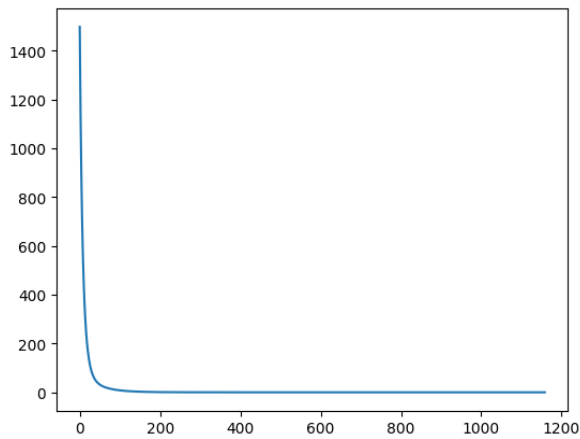
Testing data:

```
Wopt =
[1.0, 1.9999999999999999, 3.0000000000000004, 3.9999999999999999, 4.9999999999999997, 4.9999999999999996, 3.9999999999999987, 2.9999999999999996, 2.0, 1.0000000000000000]
Opt Cost (mtr) = 2.9228405874086454e-27
Opt MSE = 3.247600652676273e-28
```

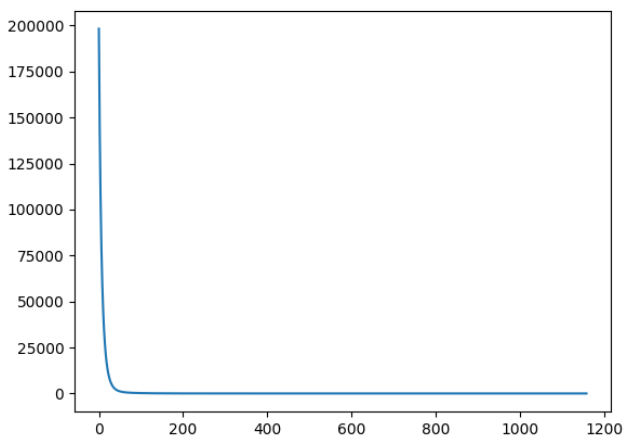
The Results: barely a difference, with the  $W_{opt}$  values nearly being the same. The testing data however, had more data and took less time to run. The training data took nearly 10x more iterations to come up with a result.

(cont.)

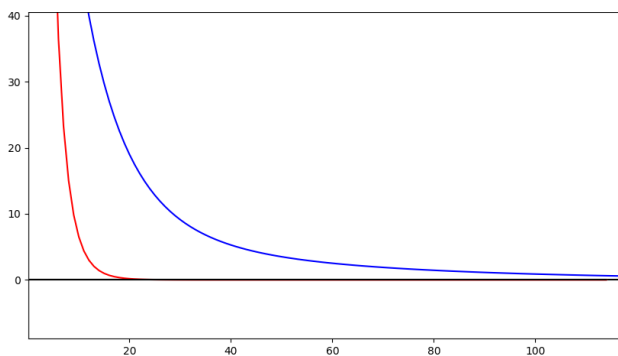
2.C – **Figure 1.** Regression objective(y) versus GD iteration (x). As the iteration progresses along the x axis, the cost metric reduces – which is expected since we want the cost to be minimal.



2.D – **Figure 2.** Euclidean norm (y) versus the iteration (x). Total # of iterations = 1158 on training data:

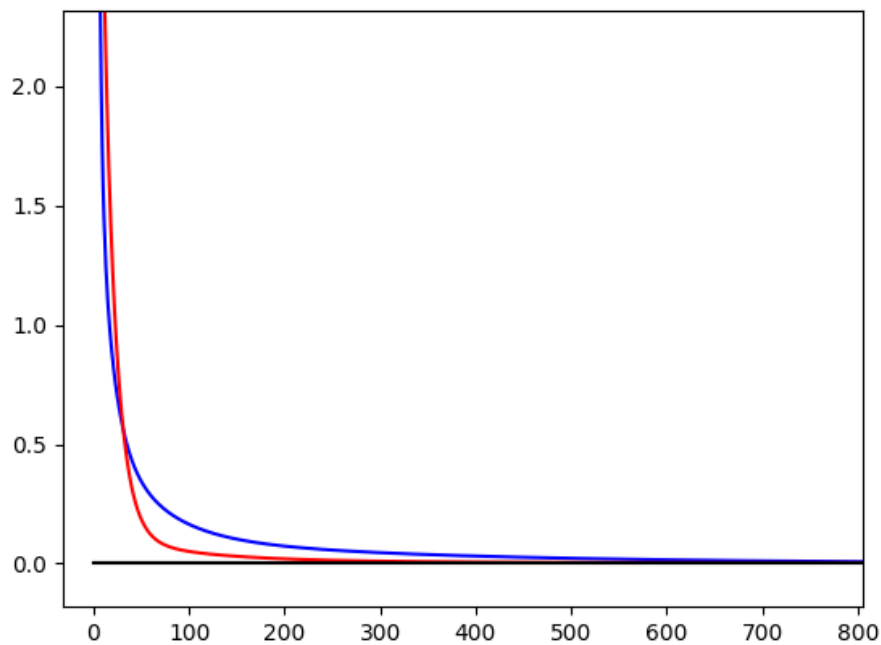


2.E – **Figure 3.** The metric lines (black) for both testing and training overlap since they are so small. Plotted are the MSE values for training (blue), testing, (red), and the benchmark MSE values (black) calculated by CF\_SOLVER. Both MSE's converge close enough to the benchmarks to consider calculation successful.



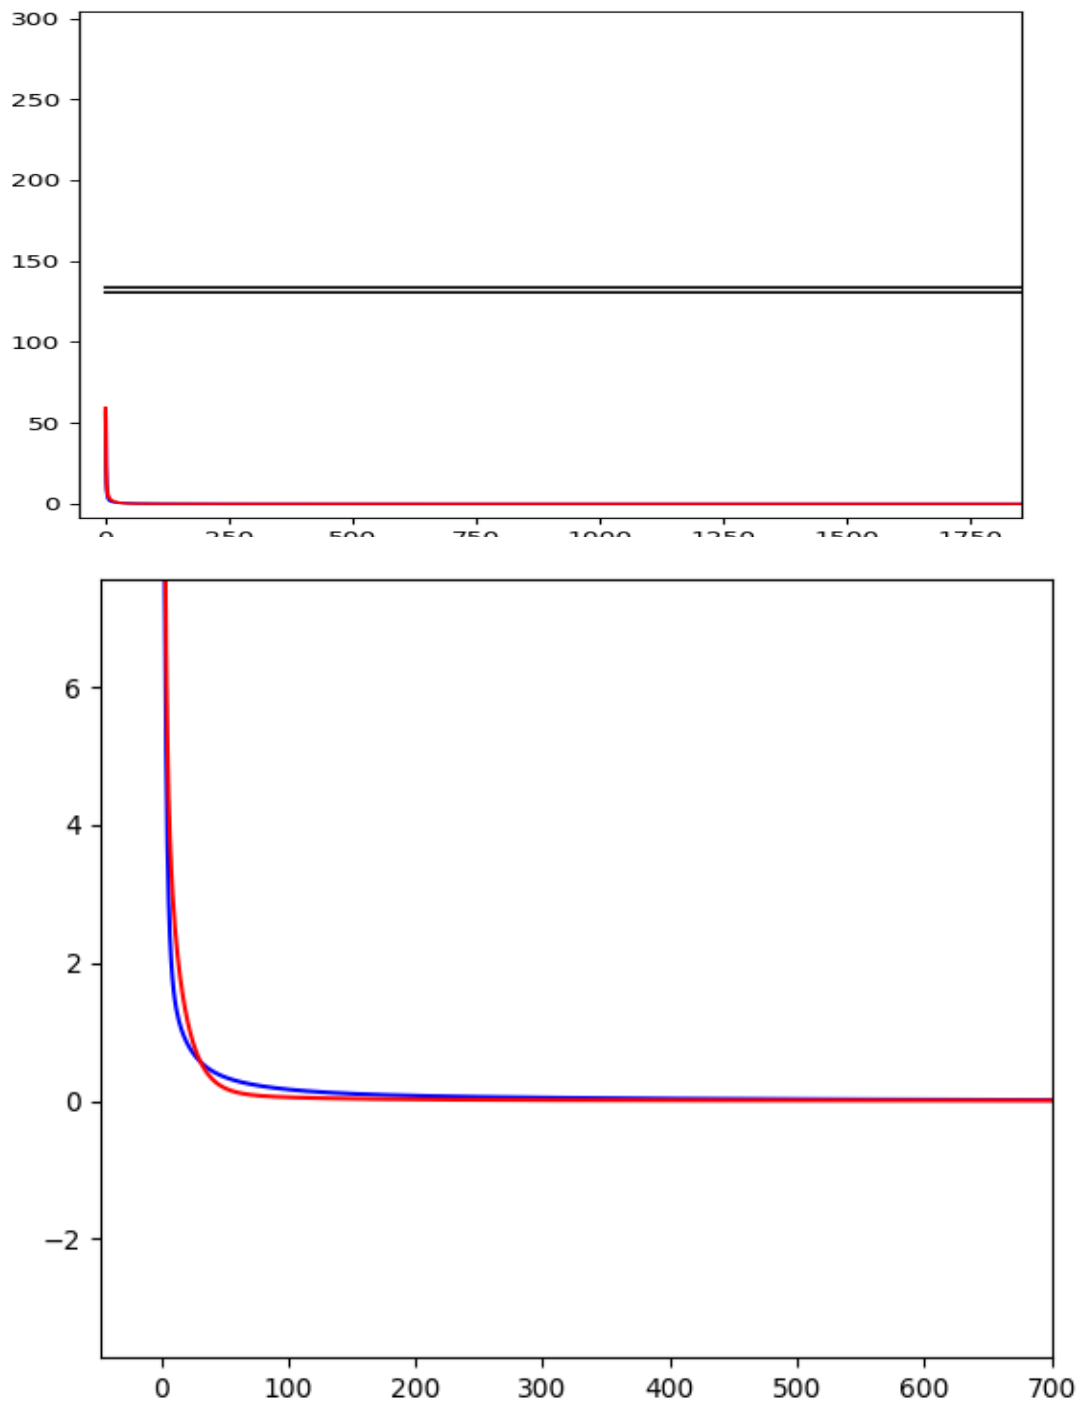
2.F – The benchmarks were meaningful since they predict a successful linear regression. Since the benchmark values aim to be at zero, the MSE lines must align with the benchmark as the iterations grow along the X axis. The MSE's are both aligning with the benchmark.

2.G – **Figure 4.** There isn't even 10 lines in the testing data to compare with... I still managed to run the plots. It seems that when the dataset is lower, more iterations are needed to converge to a minimum. Less data means more effort to calculate a relationship between variables. When there is more data, the linear relationship becomes more evident more quickly:



(cont.)

2.H – **Figure 5.** With  $\lambda = 2$ . There were nearly 5k iterations on the training data and over 20k iterations on the testing data. I assume that with a large enough  $\lambda$ , the data overfits or overshoots? The benchmark data is out of whack and doesn't represent a proper value at all. Affecting the regularization skews the accuracy:



### Problem 3 - Linear Regression (true function is affine + noise)

\*All of the following graphs

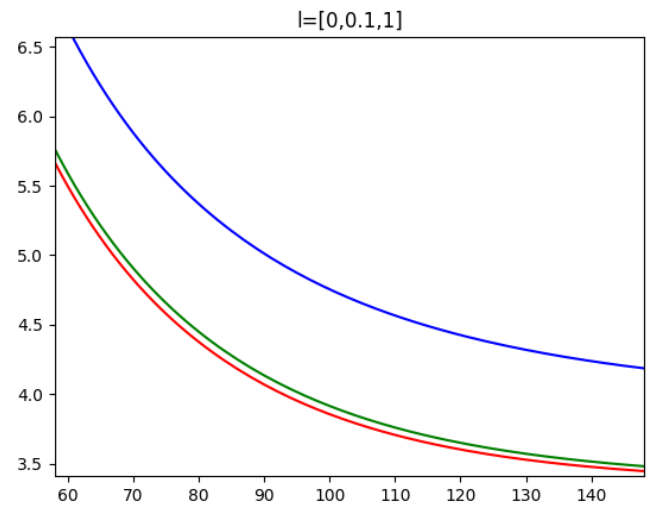
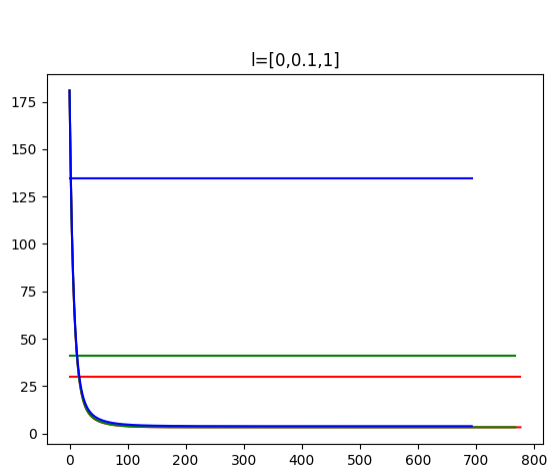
Red:  $\lambda=0$

Green:  $\lambda=0.1$

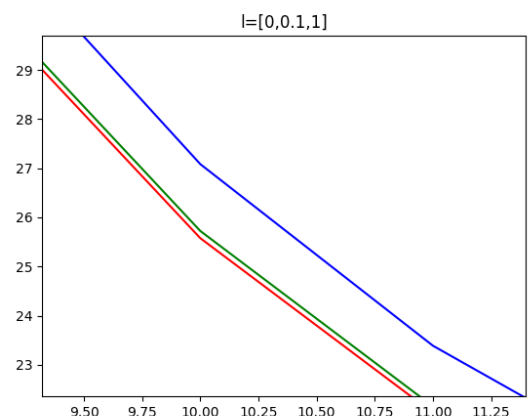
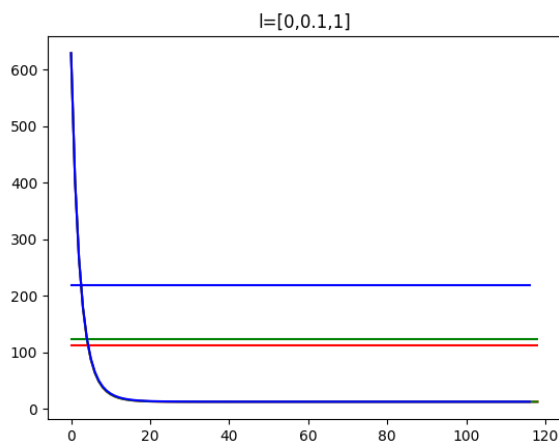
Blue:  $\lambda=1$

3.A – Check out the GitHub link : [https://github.com/ajl3545/Al\\_HW3/blob/main/problem3.py](https://github.com/ajl3545/Al_HW3/blob/main/problem3.py) or see attached files problem3.py.

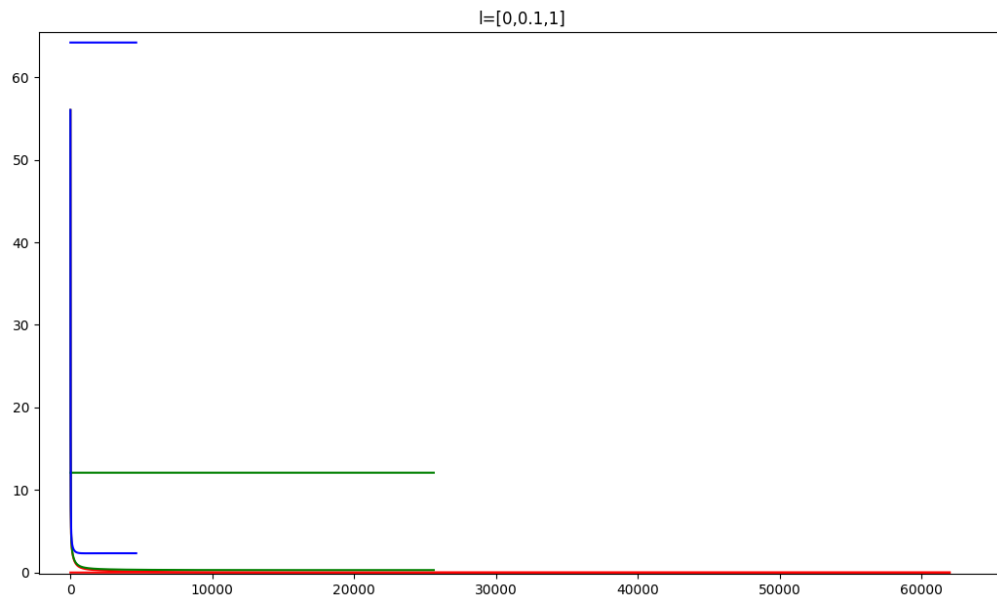
3.B – **Figure 6-7** For  $\lambda = [0, 0.1, 1]$  in the training data:



test:



3.C – **Figure 8.** For fewer data. Not sure I did this correctly



I don't know what the remaining questions for problem 3 are asking...

I tried my best. I still need to understand the concepts clearly. I can't even zoom into the graphs properly at this point.

## DATA

The following is a MSE converging on 0 error over the course of several GD iterations given a random set of parameters and training data from clean data and  $\lambda=0$  and  $\text{step} = 0.01$ :

CF\_SOLVER DATA START

wopt =

[0.9999999999999999, 2.00000000000000018, 2.9999999999999999, 4.0, 4.9999999999999999, 5.0, 3.9999999999999999, 2.9999999999999996, 1.9999999999999998, 1.00000000000000022]

Scientific notation. These values are nearly 0 (pretty much zero)

Opt Cost (mtr) = 3.3695454009416875e-28

Opt MSE = 3.7439393343796526e-29

CF\_SOLVER DATA END

GD\_SOLVER DATA START

w len = 111

iters len = 111

MSE 0 = 166.4035357461  
MSE 1 = 32.2847049312  
MSE 2 = 11.4500481246  
MSE 3 = 6.0198087551  
MSE 4 = 3.9534427416  
MSE 5 = 2.8626801898  
MSE 6 = 2.1549692721  
MSE 7 = 1.6486129376  
MSE 8 = 1.2711225328  
MSE 9 = 0.9846285416  
MSE 10 = 0.7652286020  
MSE 11 = 0.5962770496  
MSE 12 = 0.4656485482  
MSE 13 = 0.3643217362  
MSE 14 = 0.2855073769  
MSE 15 = 0.2240578514  
MSE 16 = 0.1760480677  
MSE 17 = 0.1384709846  
MSE 18 = 0.1090134325  
MSE 19 = 0.0858895519  
MSE 20 = 0.0677161522  
MSE 21 = 0.0534188664  
MSE 22 = 0.0421610999  
MSE 23 = 0.0332899662  
MSE 24 = 0.0262949577  
MSE 25 = 0.0207762204  
MSE 26 = 0.0164201081  
MSE 27 = 0.0129802817  
MSE 28 = 0.0102630512  
MSE 29 = 0.0081159775  
MSE 30 = 0.0064189887  
MSE 31 = 0.0050774403  
MSE 32 = 0.0040166850  
MSE 33 = 0.0031778171  
MSE 34 = 0.0025143317  
MSE 35 = 0.0019895003  
MSE 36 = 0.0015743059  
MSE 37 = 0.0012458175  
MSE 38 = 0.0009859092  
MSE 39 = 0.0007802506  
MSE 40 = 0.0006175097  
MSE 41 = 0.0004887245  
MSE 42 = 0.0003868063  
MSE 43 = 0.0003061475  
MSE 44 = 0.0002423117  
MSE 45 = 0.0001917891  
MSE 46 = 0.0001518022  
MSE 47 = 0.0001201534  
MSE 48 = 0.0000951038



MSE 49 = 0.0000752771  
MSE 50 = 0.0000595840  
MSE 51 = 0.0000471628  
MSE 52 = 0.0000373311  
MSE 53 = 0.0000295490  
MSE 54 = 0.0000233893  
MSE 55 = 0.0000185137  
MSE 56 = 0.0000146544  
MSE 57 = 0.0000115997  
MSE 58 = 0.0000091817  
MSE 59 = 0.0000072678  
MSE 60 = 0.0000057528  
MSE 61 = 0.0000045537  
MSE 62 = 0.0000036045  
MSE 63 = 0.0000028531  
MSE 64 = 0.0000022584  
MSE 65 = 0.0000017876  
MSE 66 = 0.0000014150  
MSE 67 = 0.0000011201  
MSE 68 = 0.0000008866  
MSE 69 = 0.0000007018  
MSE 70 = 0.0000005555  
MSE 71 = 0.0000004397  
MSE 72 = 0.0000003481  
MSE 73 = 0.0000002755  
MSE 74 = 0.0000002181  
MSE 75 = 0.0000001726  
MSE 76 = 0.0000001366  
MSE 77 = 0.0000001082  
MSE 78 = 0.0000000856  
MSE 79 = 0.0000000678  
MSE 80 = 0.0000000536  
MSE 81 = 0.0000000425  
MSE 82 = 0.0000000336  
MSE 83 = 0.0000000266  
MSE 84 = 0.0000000211  
MSE 85 = 0.0000000167  
MSE 86 = 0.0000000132  
MSE 87 = 0.0000000104  
MSE 88 = 0.0000000083  
MSE 89 = 0.0000000065  
MSE 90 = 0.0000000052  
MSE 91 = 0.0000000041  
MSE 92 = 0.0000000032  
MSE 93 = 0.0000000026  
MSE 94 = 0.0000000020  
MSE 95 = 0.0000000016  
MSE 96 = 0.0000000013  
MSE 97 = 0.0000000010  
MSE 98 = 0.0000000008  
MSE 99 = 0.0000000006  
MSE 100 = 0.0000000005  
MSE 101 = 0.0000000004  
MSE 102 = 0.0000000003  
MSE 103 = 0.0000000002  
MSE 104 = 0.0000000002  
MSE 105 = 0.0000000002  
MSE 106 = 0.0000000001  
MSE 107 = 0.0000000001  
MSE 108 = 0.0000000001  
MSE 109 = 0.0000000001  
MSE 110 = 0.0000000000  
GD\_SOLVER DATA END

GRAPH on NEXT PAGE

